



DevOps Shack

166 DataSources for Grafana Explained

Data sources in Grafana are systems or services that provide metrics, logs, or other data for visualization and analysis. Grafana connects to various data sources, such as databases, cloud services, and monitoring platforms, allowing users to create real-time dashboards and gain insights from multiple sources in a single interface.

- Diverse data sources are essential for tracking, analyzing, and managing data.
- Includes databases, cloud platforms, analytics, IoT systems, and monitoring tools.
- Covers 166 data sources like **Google Analytics**, **Amazon Redshift**, and **MongoDB**.
- Features advanced tools like **Prometheus**, **Jaeger**, and **InfluxDB**.
- Customer engagement platforms like **Zendesk** and **Salesforce** are also included.
- Each source is detailed with key features and real-world use cases.

Google Analytics	Adobe Analytics	Aggregations.io	Akvorado	Alertmanager	Atility plugin for Clickhouse	Amazon Athena	Amazon Aurora	Amazon Managed Service for Prometheus
Amazon Redshift	Amazon Timestream	Anodot Datasource	Apache Cassandra	Apache IoTDB	AppDynamics	AstraDB	Atlassian Statuspage	AWS IoT SiteWise
Autom	Azure Cosmos DB	Azure Data Explorer Datasources	Azure DevOps	Azure Monitor	Business Input	Business News	Business Satellite	Catchpoint
Chaos Mesh	Chickadee data source for Chickadee Cloud & MSP	ClickHouse	Cloudflare	CloudWatch	CloudWatch Alarm	CnosDB	CockroachDB	Cognite Data Fusion
CompareQueries	Cribl Search	CSV	Databricks	Datadog	Destiny Datasource	DeviceHive (Angular)	Discourse Datasource	Drone
Druid	DynamoDB	Dynatrace	Elasticsearch	Factory Historian Datasource	Falcon LogScale	Finnhub	Firestore	FlightSQL
Flowhook Webhooks Datasource	FROST SensorThings API Plugin	GCP Status	GitHub	GitLab	Google BigQuery	Google Cloud Logging	Google Cloud Monitoring	Google Cloud Trace
Google Sheets	GoogleCalendar (Angular)	Grafana Pyroscope	Graphite	GraphQL Data Source (Angular)	Groonga	hackone-datasource	Haystack	Hetzner Cloud
Highlight.js	Honeycomb	Hubble	IBM IBM Security QRadar	IBM IBM Security QRadar Suite	Infinity	InfluxDB	Instana	ITRS Group Observ
Jaeger	Jira	JSON	JSON API	JSON Fetcher	Kafka	kdb+	Loki	Locker
Materialize	MaxCompute	Microsoft SQL Server	Mock	MonetDB	MongoDB	MQTT	MySQL	Neo4j Datasource
netdata	Nestly	New Relic	Node Graph API	OData	openGemini	openhistorian (Angular)	OpenSearch	OpenTSDB
Optimiz-SaaSOne Plugin	Oracle Cloud Infrastructure Logs	Oracle Cloud Infrastructure Metrics	Oracle Database	Orbit	OSISOB-Pi	PagerDuty	Parca	Parseable
Pike Grafana Datasource Plugin	PHP	PostgreSQL	PostgreSQL	Prometheus	Prometheus AlertManager Datasource	Propel	Proton	QuestDB
Quickwit	Rabbitmq	Redis	Riemann streams	RunReveal	Salesforce	SAP HANA	Scanner	Sentry
ServiceNow	ServiceNow Cloud Observability	Shoreline	simple grpc datasource	skyWalking	Snowflake	SPARQL	Splunk	Splunk Infrastructure Monitoring
SQLite	Sofyze Datasource	Strava	Streamr	Sumo Logic	Sun and Moon	SurrealDB	SurrealDB	TDengine Datasource
Tempo	TestData	Thruk	Tokio Console	Timo	Vertica	Wavefront	WebSocket API	Wild GraphQL Data Source

1. Google Analytics

Overview:

Google Analytics is a web analytics service that tracks and reports website traffic. It provides detailed insights into visitor behavior, acquisition channels, and conversions, making it essential for digital marketing strategies.

Key Features:

- Real-time traffic analysis
- User demographics and behavior tracking
- Conversion goals and eCommerce tracking
- Source/Medium reporting

Example Use Case:

A retail website uses Google Analytics to monitor visitor traffic and track conversion rates from different marketing campaigns (e.g., social media ads).

2. Adobe Analytics

Overview:

Adobe Analytics is an enterprise-level analytics solution that offers in-depth customer journey insights. It provides tools for real-time tracking, segmentation, and advanced analysis.

Key Features:

- Multichannel data collection
- AI-powered predictive analytics
- Pathing analysis to track customer journeys
- Segmentation for targeted marketing

Example Use Case:

A media company uses Adobe Analytics to track user engagement across their mobile app and website, optimizing content to boost subscriptions.

3. Aggregations.io

Overview:

Aggregations.io is a high-performance tool designed for real-time event aggregation. It allows you to efficiently analyze and process large volumes of event data.

Key Features:

- Real-time data aggregation
- Scalable event processing
- Integration with other data lakes

Example Use Case:

A financial firm uses Aggregations.io to process real-time market data, enabling them to make quicker decisions on stock trades.

4. Akvorado**Overview:**

Akvorado is a network flow collector and visualization tool that helps organizations track data flows across their network infrastructure.

Key Features:

- Real-time flow data collection
- Visual flow tracking
- Security and performance monitoring

Example Use Case:

An IT department uses Akvorado to monitor traffic between data centers and detect abnormal patterns that might indicate security threats.

5. Alertmanager**Overview:**

Alertmanager is a tool used with Prometheus to handle alerts. It routes, groups, silences, and deduplicates alerts.

Key Features:

- Alert grouping and routing
- Deduplication of alerts
- Customizable alerting policies

Example Use Case:

A DevOps team uses Alertmanager to manage alerts from Prometheus and ensure only critical issues are escalated to on-call engineers.

6. Altinity Plugin for ClickHouse**Overview:**

Altinity provides enterprise-grade plugins for ClickHouse, a columnar database designed for analytical queries on large datasets.

Key Features:

- Improved ClickHouse performance
- Enhanced security features

- Integration with BI tools like Grafana

Example Use Case:

A marketing company uses ClickHouse with the Altinity plugin to analyze customer behavior across multiple channels, providing real-time insights.

7. Amazon Athena

Overview:

Amazon Athena is a serverless query service that makes it easy to analyze data stored in Amazon S3 using standard SQL.

Key Features:

- Serverless querying of S3 data
- SQL support
- Pay-per-query pricing model

Example Use Case:

A retail business uses Athena to run SQL queries on their S3-stored eCommerce data to analyze sales patterns without setting up complex infrastructure.

8. Amazon Aurora

Overview:

Amazon Aurora is a fully managed MySQL and PostgreSQL-compatible relational database. It is designed for high availability and scalability.

Key Features:

- High-performance with scalability
- Fully managed and automated backups
- Multi-region replication

Example Use Case:

An online gaming company uses Aurora to store player profiles, ensuring fast reads and writes even during peak usage times.

9. Amazon Managed Service for Prometheus

Overview:

A fully managed Prometheus-compatible monitoring service, Amazon Managed Service for Prometheus allows users to monitor containerized applications at scale.

Key Features:

- Fully managed service with Prometheus compatibility

- Scalable for large containerized environments
- Integration with AWS services

Example Use Case:

A SaaS provider monitors the performance and health of their Kubernetes clusters using this managed service to reduce operational overhead.

10. Amazon Redshift

Overview:

Amazon Redshift is a fast, scalable data warehouse that allows for large-scale data analysis.

Key Features:

- Columnar storage for performance optimization
- Massively parallel processing (MPP)
- Integration with S3 for data loading

Example Use Case:

A logistics company uses Redshift to analyze large datasets, helping optimize shipping routes based on historical data.

11. Amazon Timestream

Overview:

Amazon Timestream is a time-series database optimized for IoT and operational applications that process time-based data.

Key Features:

- Serverless time-series database
- Built-in analytics functions
- High throughput for real-time data

Example Use Case:

A smart home company uses Timestream to monitor sensor data from devices in real-time, helping to optimize energy consumption.

12. Anodot Datasource

Overview:

Anodot is an AI-driven analytics platform that helps with anomaly detection in business metrics by collecting and monitoring data streams in real-time.

Key Features:

- AI-based anomaly detection

- Real-time data monitoring
- Root cause analysis for anomalies

Example Use Case:

An eCommerce company uses Anodot to monitor sales metrics and detect any sudden drops in transactions due to potential technical issues.

13. Apache Cassandra

Overview:

Apache Cassandra is a distributed NoSQL database designed for handling large amounts of data with no single point of failure, making it ideal for high-availability applications.

Key Features:

- Linear scalability
- High availability and fault tolerance
- Tunable consistency levels

Example Use Case:

A social media platform uses Cassandra to store and retrieve user-generated content across multiple data centers, ensuring high availability.

14. Apache IoTDB

Overview:

Apache IoTDB is a time-series database specifically designed for storing IoT data. It supports high write throughput and compression for large-scale IoT applications.

Key Features:

- Optimized for time-series data
- Efficient data compression
- High write throughput

Example Use Case:

A manufacturing company uses Apache IoTDB to monitor machinery health by collecting sensor data in real-time, reducing downtime through predictive maintenance.

15. AppDynamics

Overview:

AppDynamics is an application performance monitoring (APM) tool that provides real-time performance insights and analytics for your applications and infrastructure.

Key Features:

- End-to-end transaction tracing
- Real-time performance monitoring
- AI-driven insights and anomaly detection

Example Use Case:

A banking institution uses AppDynamics to monitor its online banking system, ensuring that transactions are processed without delays.

16. AstraDB

Overview:

AstraDB is a serverless NoSQL database built on Apache Cassandra, designed to simplify the management of large datasets while providing robust scalability and high availability.

Key Features:

- Built on Apache Cassandra
- Serverless architecture for scalability
- Simplified database management

Example Use Case:

A fintech company uses AstraDB to store and retrieve transactional data from millions of users, benefiting from its scalability and reliability.

17. Atlassian Statuspage

Overview:

Atlassian Statuspage is a communication tool that allows businesses to inform customers about the status of their services, incidents, and scheduled maintenance.

Key Features:

- Real-time status updates
- Incident communication and subscription management
- Customizable status pages

Example Use Case:

A cloud service provider uses Statuspage to keep their customers informed of any service outages or planned maintenance, reducing the volume of customer support queries during incidents.

18. AWS IoT SiteWise

Overview:

AWS IoT SiteWise is a managed service that makes it easy to collect, organize, and analyze data from industrial equipment at scale.

Key Features:

- Collect and monitor IoT data from industrial devices
- Build and use asset models for real-time insights
- Real-time visualizations and analysis

Example Use Case:

A factory collects and monitors data from machinery to optimize operations and perform predictive maintenance, reducing costly unplanned downtime.

19. Axiom**Overview:**

Axiom is a cloud-native platform designed to collect and process event data at scale. It helps analyze log and event data in real-time for actionable insights.

Key Features:

- Real-time data ingestion and analytics
- Scalable, cloud-native architecture
- Log data visualization

Example Use Case:

A security operations center (SOC) uses Axiom to collect and analyze log data from thousands of systems, helping detect and respond to potential security incidents.

20. Azure Cosmos DB**Overview:**

Azure Cosmos DB is a fully managed NoSQL database service for building globally distributed, scalable applications with support for multiple data models.

Key Features:

- Multi-model (document, graph, key-value, column-family)
- Global distribution with low latency
- SLA-backed availability, throughput, and consistency

Example Use Case:

An online retail business uses Cosmos DB to store and manage product catalogs and customer data across multiple regions, ensuring low-latency access for users worldwide.

21. Azure Data Explorer Datasource

Overview:

Azure Data Explorer is a fully managed data analytics service optimized for querying large volumes of log and telemetry data in real-time.

Key Features:

- Real-time analytics for log and telemetry data
- Fast and scalable queries on large datasets
- Integrated with Azure services

Example Use Case:

A cloud monitoring company uses Azure Data Explorer to process and analyze log data from thousands of cloud instances, providing actionable insights to their operations team.

22. Azure DevOps

Overview:

Azure DevOps is a suite of development tools for planning, developing, testing, and deploying applications. It includes services like Azure Boards, Repos, Pipelines, Test Plans, and Artifacts.

Key Features:

- CI/CD pipelines for automated deployment
- Source code management with Git
- Integrated testing and artifact management

Example Use Case:

A software development team uses Azure DevOps Pipelines to automate the testing and deployment of their application, ensuring faster release cycles.

23. Azure Monitor

Overview:

Azure Monitor is a comprehensive monitoring service for collecting, analyzing, and acting on telemetry data from Azure resources and applications.

Key Features:

- Metrics and logs collection for Azure resources
- Integrated with alerting and automated actions
- Visualizations through Azure dashboards and Power BI

Example Use Case:

An eCommerce platform uses Azure Monitor to track the performance of its virtual machines, databases, and web apps, ensuring system reliability and performance.

24. Business Input

Overview:

Business Input refers to internal or external data collected from business operations or activities, such as customer feedback, sales data, or market research.

Key Features:

- Custom data collection from business operations
- Integration with BI tools for analysis
- Actionable insights for decision-making

Example Use Case:

A retail chain collects feedback from customers after purchases to adjust their product offerings and improve customer satisfaction.

25. Business News

Overview:

Business News data sources provide up-to-date information about market trends, financial updates, and other economic indicators that affect businesses.

Key Features:

- Real-time market and financial news
- Access to industry-specific reports
- Market trend analysis

Example Use Case:

A stock trader uses Business News to stay updated on changes in the global financial market, allowing them to make informed trading decisions.

26. Business Satellite

Overview:

Business Satellite data sources provide satellite-based insights, such as geospatial data, that are critical for industries like agriculture, logistics, and disaster management.

Key Features:

- Geospatial data collection via satellite imagery
- Integration with GIS platforms
- Real-time monitoring and analytics

Example Use Case:

An agricultural company uses Business Satellite data to monitor crop health and predict yield, helping optimize resource allocation.

27. Catchpoint

Overview:

Catchpoint is a digital experience monitoring (DEM) platform that provides visibility into the performance of websites, APIs, and applications from a global perspective.

Key Features:

- Synthetic and real user monitoring
- Global endpoint testing
- Root cause analysis of performance issues

Example Use Case:

An online streaming service uses Catchpoint to monitor the performance of their platform across different regions, ensuring a smooth user experience regardless of location.

28. Chaos Mesh

Overview:

Chaos Mesh is a cloud-native Chaos Engineering platform that helps you inject faults into Kubernetes applications to test their resilience.

Key Features:

- Fault injection for various failure scenarios
- Integrates with Kubernetes environments
- Resilience testing for cloud-native applications

Example Use Case:

A DevOps team uses Chaos Mesh to test the fault tolerance of their microservices architecture, ensuring the application can handle unexpected disruptions in production.

29. Checkmk data source for Checkmk Cloud & MSP

Overview:

Checkmk is a comprehensive monitoring tool for infrastructure and applications, and the data source for Checkmk Cloud & MSP enables multi-tenant monitoring for managed service providers.

Key Features:

- Real-time infrastructure and application monitoring
- Multi-tenant monitoring support
- Automated alerting and reporting

Example Use Case:

A managed service provider uses Checkmk Cloud to monitor and manage IT infrastructure across multiple customer environments from a single interface.

30. ClickHouse**Overview:**

ClickHouse is a fast, open-source columnar database management system designed for real-time data analytics.

Key Features:

- High-speed data ingestion and query execution
- Columnar storage format for optimized analytics
- Supports petabyte-scale data storage

Example Use Case:

A financial institution uses ClickHouse to analyze large datasets containing stock market transactions and generate real-time trading reports.

31. Cloudflare**Overview:**

Cloudflare is a web infrastructure and security company providing content delivery network (CDN) services, DDoS mitigation, Internet security, and distributed domain name server services.

Key Features:

- Web security and DDoS protection
- CDN for faster content delivery
- DNS and SSL/TLS management

Example Use Case:

An online retailer uses Cloudflare's CDN and security services to ensure their website loads quickly and remains secure during traffic surges.

32. CloudWatch**Overview:**

Amazon CloudWatch is a monitoring and observability service for AWS cloud resources and applications, providing actionable insights based on metrics and logs.

Key Features:

- Real-time monitoring and alerting
- Custom dashboards for visualization

- Log aggregation and search

Example Use Case:

A SaaS provider uses CloudWatch to monitor the performance of their AWS infrastructure, setting up alerts for CPU usage spikes and potential system failures.

33. CloudWatch Alarm

Overview:

CloudWatch Alarms are part of Amazon CloudWatch and allow users to set alarms that trigger notifications or automated actions based on monitored metrics.

Key Features:

- Set custom thresholds for alarms
- Trigger AWS Lambda functions or EC2 actions
- Integration with SNS for alerting

Example Use Case:

A DevOps team uses CloudWatch Alarms to automatically scale their EC2 instances when CPU usage exceeds a specified threshold, ensuring their application handles traffic spikes.

34. CnosDB

Overview:

CnosDB is a time-series database optimized for high-velocity data ingestion, storage, and querying, often used for IoT and telemetry applications.

Key Features:

- Time-series data storage and analysis
- High-velocity data ingestion
- Optimized for IoT applications

Example Use Case:

An IoT company uses CnosDB to collect and store sensor data from smart devices, enabling real-time monitoring and analytics of environmental conditions.

35. CockroachDB

Overview:

CockroachDB is a distributed SQL database that provides strong consistency, fault tolerance, and horizontal scalability, ideal for globally distributed applications.

Key Features:

- Strong consistency and ACID transactions

- Horizontal scalability across regions
- Fault-tolerant architecture

Example Use Case:

A global eCommerce platform uses CockroachDB to ensure that customer transactions are consistent and fault-tolerant across multiple regions, minimizing downtime during outages.

36. Cognite Data Fusion

Overview:

Cognite Data Fusion is an industrial data operations platform that enables the collection, contextualization, and analysis of data from various industrial sources, such as sensors and machinery.

Key Features:

- Real-time industrial data collection
- Data contextualization for analysis
- Integration with AI and machine learning tools

Example Use Case:

An oil and gas company uses Cognite Data Fusion to analyze sensor data from oil rigs, helping optimize production and reduce equipment downtime.

37. CompareQueries

Overview:

CompareQueries is a query performance analysis tool that helps developers compare SQL queries across different databases and versions to optimize performance.

Key Features:

- Query performance comparison
- Supports multiple database engines
- Detailed performance metrics and analysis

Example Use Case:

A data analyst uses CompareQueries to benchmark different SQL queries on MySQL and PostgreSQL databases to identify performance bottlenecks and optimize reporting queries.

38. Cribl Search

Overview:

Cribl Search is a data search and streaming platform that allows organizations to search through event data, logs, and metrics in real-time from multiple sources without needing to store the data first.

Key Features:

- Search through event data in real-time
- Stream data without the need for storage
- Supports multiple data sources

Example Use Case:

An IT team uses Cribl Search to investigate a security breach by searching through logs from multiple systems in real-time without having to store the logs first.

39. CSV**Overview:**

CSV (Comma Separated Values) is a widely-used format for storing tabular data in a plain-text file, where each line represents a row, and each value is separated by a comma.

Key Features:

- Simple text-based data format
- Easily parsed by most programming languages
- Lightweight and human-readable

Example Use Case:

A sales team exports monthly sales data from their CRM system into a CSV file for quick reporting and analysis in Excel.

40. Databricks**Overview:**

Databricks is a unified analytics platform for big data and machine learning, providing a collaborative environment for data engineers, scientists, and business analysts.

Key Features:

- Unified analytics for big data and AI
- Collaborative environment with notebooks
- Supports Apache Spark for large-scale data processing

Example Use Case:

A pharmaceutical company uses Databricks to analyze large datasets from clinical trials, applying machine learning models to predict drug efficacy.

41. Datadog

Overview:

Datadog is a cloud monitoring and observability platform that provides monitoring, alerting, and visualization of metrics from servers, databases, tools, and services.

Key Features:

- Monitoring for infrastructure, applications, and logs
- Customizable dashboards and alerts
- Integration with over 400 services and tools

Example Use Case:

A cloud-native company uses Datadog to monitor their AWS infrastructure, creating custom alerts for issues like high CPU utilization and low disk space.

42. Destiny Datasource

Overview:

Destiny is a data integration and management platform that enables organizations to collect, store, and analyze data from multiple sources in one place.

Key Features:

- Data integration from multiple sources
- Centralized storage for data analysis
- Scalable architecture for large datasets

Example Use Case:

A logistics company uses Destiny to integrate data from their fleet management system, warehouse inventory, and customer orders, enabling them to optimize delivery routes and inventory levels.

43. DeviceHive

Overview:

DeviceHive is an open-source IoT platform that provides device management, data collection, and real-time control for IoT applications.

Key Features:

- Real-time device management and control
- Supports various IoT protocols (MQTT, CoAP, etc.)
- Data collection and analytics

Example Use Case:

A smart city project uses DeviceHive to manage a network of connected streetlights, collecting data on energy consumption and controlling lighting based on real-time conditions.

44. Discourse Datasource

Overview:

Discourse is a modern forum platform designed for online discussions and community engagement. The Discourse Datasource allows you to analyze user engagement and forum activity.

Key Features:

- User engagement tracking
- Real-time discussion data collection
- Integration with community management tools

Example Use Case:

A company with a large online forum uses Discourse to monitor user activity, tracking engagement levels and identifying the most popular discussion topics.

45. Drone

Overview:

Drone is a continuous integration and delivery (CI/CD) platform that automates the testing and deployment of applications, providing a streamlined pipeline for DevOps teams.

Key Features:

- Automated CI/CD pipelines
- Integration with Git repositories
- Docker support for containerized builds

Example Use Case:

A software development team uses Drone to automate the testing and deployment of their microservices, reducing the time and effort needed to push updates to production.

46. Druid

Overview:

Apache Druid is a high-performance, real-time analytics database designed for fast query execution on large datasets. It's often used for real-time dashboards and event-driven applications.

Key Features:

- Real-time data ingestion and querying
- Fast query execution on large datasets
- Ideal for time-series and event-based data

Example Use Case:

A media company uses Druid to analyze user engagement on their streaming platform in real-time, helping optimize content recommendations.

47. DynamoDB

Overview:

Amazon DynamoDB is a fully managed NoSQL database service that provides fast, predictable performance and scalability for internet-scale applications.

Key Features:

- Fully managed NoSQL database
- Millisecond response times at any scale
- Global tables for multi-region replication

Example Use Case:

An online gaming company uses DynamoDB to store player data in a globally distributed fashion, ensuring fast response times and high availability for players worldwide.

48. Dynatrace

Overview:

Dynatrace is a software intelligence platform that provides observability into applications, infrastructure, and user experience through AI-powered analytics.

Key Features:

- AI-driven insights and anomaly detection
- Full-stack observability across applications and infrastructure
- Real-user monitoring for digital experience

Example Use Case:

A financial services company uses Dynatrace to monitor the performance of their customer-facing applications, ensuring that any issues are automatically detected and resolved.

49. Elasticsearch

Overview:

Elasticsearch is a distributed, RESTful search and analytics engine built on top of Apache Lucene. It is used for full-text search, log analytics, and real-time data visualization.

Key Features:

- Full-text search engine with real-time indexing
- Scalable and distributed architecture
- Real-time analytics and dashboarding (with Kibana)

Example Use Case:

A cybersecurity company uses Elasticsearch to store and search through millions of logs in real-time, helping them detect potential security breaches.

50. Factory Historian Datasource**Overview:**

Factory Historian is an industrial data historian software that captures and stores high-frequency data from sensors and machinery for analysis.

Key Features:

- High-frequency industrial data collection
- Time-series data storage and visualization
- Integration with SCADA and other industrial systems

Example Use Case:

A factory uses Factory Historian to capture sensor data from their production line, enabling real-time monitoring of machine performance and predictive maintenance.

51. Falcon LogScale**Overview:**

Falcon LogScale is a scalable log management solution that enables organizations to store, search, and analyze large volumes of log data in real-time.

Key Features:

- Scalable log management and analytics
- Real-time log search and visualization
- Integration with SIEM and security tools

Example Use Case:

An IT security team uses Falcon LogScale to manage logs from thousands of servers, allowing them to quickly detect and respond to potential security threats.

52. Finnhub**Overview:**

Finnhub is a financial data API providing real-time stock, forex, and cryptocurrency market data, along with financial news and economic indicators.

Key Features:

- Real-time financial data for stocks, forex, and crypto
- News and economic indicators

- Company profiles and fundamentals

Example Use Case:

A stock trading platform uses Finnhub to display real-time stock prices and market news, helping traders make informed decisions.

53. Firestore

Overview:

Firestore is a cloud-hosted NoSQL database from Google Firebase that provides real-time data syncing and offline capabilities for web and mobile applications.

Key Features:

- Real-time data synchronization
- Offline support for mobile apps
- Scalable, fully managed NoSQL database

Example Use Case:

A mobile chat application uses Firestore to store and sync messages in real-time, even when users are offline, ensuring that all messages are available once they reconnect.

54. FlightSQL

Overview:

FlightSQL is a protocol for interacting with large datasets using SQL queries, optimized for data transfer in distributed systems and large-scale analytics.

Key Features:

- SQL-based querying protocol
- Optimized for high-speed data transfer
- Ideal for large-scale distributed systems

Example Use Case:

A big data analytics company uses FlightSQL to perform distributed queries on large datasets, providing fast results even for complex queries.

55. Flowhook Webhooks Datasource

Overview:

Flowhook is a webhook management platform that allows you to create, manage, and monitor webhooks for real-time data integration and event-driven workflows.

Key Features:

- Webhook management and monitoring

- Real-time data integration with other services
- Customizable event-driven workflows

Example Use Case:

A cloud-based CRM system uses Flowhook to trigger webhooks when a new lead is created, automatically sending the data to their marketing automation platform for follow-up actions.

56. FROST SensorThings API Plugin

Overview:

FROST SensorThings is an API standard for IoT devices to communicate sensor data in a consistent format, making it easy to collect, manage, and analyze IoT data.

Key Features:

- Standardized API for IoT sensor data
- Real-time data collection and analysis
- Supports multiple IoT protocols

Example Use Case:

A smart city project uses FROST SensorThings to collect environmental data from air quality sensors, providing real-time insights into pollution levels.

57. GCP Status

Overview:

GCP Status provides real-time updates and historical information on the status of Google Cloud Platform services, including outages and maintenance events.

Key Features:

- Real-time status updates for GCP services
- Incident tracking and resolution timelines
- Historical data on service uptime

Example Use Case:

A DevOps team uses GCP Status to monitor the availability of their Google Cloud resources and respond quickly to any outages or service disruptions.

58. GitHub

Overview:

GitHub is a web-based version control and collaboration platform for software developers, allowing teams to collaborate on code, track changes, and manage repositories.

Key Features:

- Git-based version control
- Collaborative tools for teams
- Integrated CI/CD pipelines with GitHub Actions

Example Use Case:

A software development team uses GitHub to collaborate on a large open-source project, managing code contributions from developers across the globe.

59. GitLab

Overview:

GitLab is a DevOps platform that integrates version control, CI/CD pipelines, and project management tools in one platform, enabling end-to-end software development workflows.

Key Features:

- Integrated CI/CD pipelines
- Git-based version control
- Project management and issue tracking

Example Use Case:

A DevOps team uses GitLab to manage their entire software development lifecycle, from coding to deployment, with automated testing and deployment pipelines.

60. Google BigQuery

Overview:

Google BigQuery is a fully managed, serverless data warehouse designed for large-scale analytics, allowing you to query massive datasets using standard SQL.

Key Features:

- Serverless, fully managed data warehouse
- Supports SQL queries on large datasets
- Real-time analytics and ML integration

Example Use Case:

A marketing team uses BigQuery to analyze customer data from various sources, running SQL queries to identify trends and optimize their ad campaigns.

61. Google Cloud Logging

Overview:

Google Cloud Logging is a fully managed log management service that allows you to store, search, analyze, and monitor log data from Google Cloud resources.

Key Features:

- Real-time log data collection and analysis
- Integration with Google Cloud services
- Customizable log-based metrics and alerts

Example Use Case:

A cloud-native application uses Google Cloud Logging to monitor log data from its Google Kubernetes Engine (GKE) clusters, ensuring that issues are detected and addressed quickly.

62. Google Cloud Monitoring

Overview:

Google Cloud Monitoring provides monitoring and observability services for Google Cloud resources and applications, helping users track metrics, set up dashboards, and create alerts.

Key Features:

- Monitoring for Google Cloud resources and apps
- Customizable dashboards and alerts
- Integration with logs and events

Example Use Case:

A SaaS provider uses Google Cloud Monitoring to track the performance of its virtual machines and databases, creating alerts for high CPU usage and memory consumption.

63. Google Cloud Trace

Overview:

Google Cloud Trace is a distributed tracing service for analyzing the performance of applications by tracking latency across microservices and requests.

Key Features:

- Distributed tracing for microservices
- Real-time performance analysis
- Integration with Google Cloud services

Example Use Case:

A microservices-based application uses Google Cloud Trace to analyze the performance of different services, identifying bottlenecks and optimizing response times.

64. Google Sheets

Overview:

Google Sheets is a cloud-based spreadsheet application that allows for real-time collaboration and integration with other Google services like Google Drive and Google Apps Script.

Key Features:

- Real-time collaboration on spreadsheets
- Cloud-based storage and access
- Integration with Google Apps Script for automation

Example Use Case:

A small business uses Google Sheets to track their monthly expenses and sales data, collaborating with their accountant in real-time to manage finances.

65. GoogleCalendar

Overview:

Google Calendar is a cloud-based scheduling tool that allows users to create, share, and manage events and appointments, with integration into other Google services.

Key Features:

- Event scheduling and reminders
- Integration with Gmail, Google Drive, and other apps
- Real-time collaboration on shared calendars

Example Use Case:

A project manager uses Google Calendar to schedule meetings and deadlines for the development team, ensuring everyone is aligned on timelines.

66. Grafana Pyroscope

Overview:

Pyroscope is an open-source continuous profiling platform that allows developers to profile their applications and monitor CPU, memory, and I/O usage over time.

Key Features:

- Continuous profiling of application performance
- Real-time monitoring of CPU, memory, and I/O usage
- Integration with Grafana for visualization

Example Use Case:

A backend development team uses Pyroscope to identify performance bottlenecks in their API, optimizing memory usage and reducing latency.

67. Graphite

Overview:

Graphite is an open-source monitoring tool that stores and visualizes time-series data, often used for tracking server performance metrics such as CPU usage, memory, and disk I/O.

Key Features:

- Real-time collection and storage of time-series data
- Powerful querying and visualization tools
- Custom dashboards for performance monitoring

Example Use Case:

An IT team uses Graphite to monitor the health of their server infrastructure, tracking CPU and memory usage to ensure uptime and performance.

68. GraphQL Data Source

Overview:

GraphQL is a query language for APIs and a runtime for executing those queries, enabling more flexible data retrieval and manipulation compared to REST APIs.

Key Features:

- Flexible data querying and retrieval
- Type-safe API with real-time introspection
- Reduces over-fetching of data

Example Use Case:

A mobile app uses GraphQL to fetch only the necessary data for a specific screen, improving performance and reducing bandwidth usage compared to traditional REST APIs.

69. Groonga

Overview:

Groonga is an open-source full-text search engine and column store, designed for fast text searching and real-time updates.

Key Features:

- Full-text search with fast indexing
- Real-time data updates
- Column store for efficient data retrieval

Example Use Case:

A news website uses Groonga to provide fast, accurate search results for articles, allowing users to quickly find content based on keywords and topics.

70. hackerone-datasource

Overview:

HackerOne is a platform that connects businesses with ethical hackers to identify vulnerabilities in their systems. The hackerone-datasource plugin allows for tracking and analyzing vulnerabilities discovered via the platform.

Key Features:

- Vulnerability tracking and reporting
- Integration with security tools and dashboards
- Real-time updates on discovered vulnerabilities

Example Use Case:

A financial services company uses HackerOne to manage their bug bounty program, allowing ethical hackers to report vulnerabilities and helping the company prioritize and fix security issues.

71. Haystack

Overview:

Haystack is an observability platform designed to monitor and debug microservices applications, providing distributed tracing, metrics, and logs for performance analysis.

Key Features:

- Distributed tracing for microservices
- Real-time performance monitoring
- Metrics and logs integration

Example Use Case:

A microservices-based eCommerce platform uses Haystack to trace transactions across services, identifying bottlenecks and improving user experience during peak shopping times.

72. Hetzner Cloud

Overview:

Hetzner Cloud is a cloud hosting provider offering virtual private servers (VPS), dedicated servers, and storage solutions for web applications and businesses.

Key Features:

- Scalable virtual and dedicated server hosting
- High-performance infrastructure with SSD storage
- API and CLI for automation and management

Example Use Case:

A web development agency uses Hetzner Cloud to host client websites, taking advantage of its high-performance infrastructure and flexible pricing.

73. Highlight.io**Overview:**

Highlight.io is a monitoring tool that provides real-time insights into user behavior and performance on web applications, helping developers optimize UX and performance.

Key Features:

- Real-time user session tracking
- Performance monitoring and analytics
- Error tracking and resolution

Example Use Case:

A web development team uses Highlight.io to monitor user behavior on their web app, identifying slow-loading pages and improving performance.

74. Honeycomb**Overview:**

Honeycomb is a platform for observability, offering high-cardinality analysis of complex system behavior through event-based tracing, metrics, and logs.

Key Features:

- High-cardinality event-based tracing
- Real-time insights into application performance
- Integrated metrics and logs for comprehensive observability

Example Use Case:

A DevOps team uses Honeycomb to investigate and resolve performance issues in a distributed microservices application, pinpointing the root cause of latency spikes.

75. Hubble**Overview:**

Hubble is a Kubernetes-native observability platform that provides deep visibility into network traffic and security policies across your Kubernetes clusters.

Key Features:

- Network observability for Kubernetes clusters
- Real-time monitoring of security policies

- Integration with CI/CD pipelines for security

Example Use Case:

A security-conscious DevOps team uses Hubble to monitor network traffic in their Kubernetes cluster, ensuring compliance with security policies and preventing unauthorized access.

76. IBM Security QRadar**Overview:**

IBM Security QRadar is a security information and event management (SIEM) tool that helps organizations detect and respond to security threats by collecting, normalizing, and analyzing security data.

Key Features:

- SIEM platform with real-time threat detection
- Security data normalization and correlation
- Automated incident response and remediation

Example Use Case:

A large enterprise uses QRadar to collect security logs from its network devices, servers, and applications, using machine learning to detect and respond to potential cyberattacks.

77. IBM Security QRadar Suite**Overview:**

IBM Security QRadar Suite builds on the core SIEM functionality of QRadar by adding more advanced threat intelligence, analytics, and incident response features.

Key Features:

- Advanced threat detection and intelligence
- Integrated incident response workflows
- Analytics for proactive threat hunting

Example Use Case:

A security operations center (SOC) uses QRadar Suite to proactively hunt for threats across their network, using threat intelligence feeds to detect and respond to new cyber threats.

78. Infinity**Overview:**

Infinity is a data integration and orchestration platform that enables businesses to create, manage, and automate data pipelines from multiple sources into a unified view.

Key Features:

- Data integration from multiple sources

- Automated pipeline creation and management
- Real-time data processing

Example Use Case:

A retail business uses Infinity to integrate data from their eCommerce platform, CRM, and marketing tools into a unified dashboard for tracking sales performance.

79. InfluxDB**Overview:**

InfluxDB is an open-source time-series database designed for storing high-velocity data such as metrics and events, often used for monitoring and IoT applications.

Key Features:

- Optimized for time-series data
- Real-time data ingestion and querying
- Built-in support for retention policies and downsampling

Example Use Case:

A smart factory uses InfluxDB to store sensor data from machines on the production line, enabling real-time monitoring and predictive maintenance.

80. Instana**Overview:**

Instana is an application performance monitoring (APM) platform that provides real-time insights into the performance of applications, infrastructure, and microservices.

Key Features:

- Automated discovery and monitoring of applications
- Real-time performance insights and alerts
- Full-stack observability for microservices

Example Use Case:

A microservices-based application uses Instana to monitor the performance of each service, ensuring any performance degradation is detected and addressed immediately.

81. ITRS Group Obcerv**Overview:**

ITRS Obcerv is a data observability platform that captures and stores data from IT infrastructure, applications, and network devices, enabling real-time monitoring and historical analysis.

Key Features:

- Real-time data collection from IT infrastructure
- Historical analysis and reporting
- Scalable architecture for large environments

Example Use Case:

A financial services company uses ITRS Obcerv to monitor the performance of their trading platform, ensuring uptime and performance during high-traffic trading hours.

82. Jaeger**Overview:**

Jaeger is an open-source distributed tracing system designed for monitoring and troubleshooting microservices-based architectures, helping developers track the flow of requests across services.

Key Features:

- Distributed tracing for microservices
- Performance monitoring and troubleshooting
- Integration with Kubernetes and other orchestration tools

Example Use Case:

A DevOps team uses Jaeger to trace requests across their microservices architecture, identifying bottlenecks in the system and optimizing service performance.

83. Jira**Overview:**

Jira is a project management tool used by software development teams for tracking issues, managing workflows, and collaborating on code. It integrates with many CI/CD tools for DevOps automation.

Key Features:

- Issue tracking and project management
- Customizable workflows and automation
- Integration with CI/CD tools

Example Use Case:

A software development team uses Jira to track bugs and feature requests, ensuring that all tasks are organized and prioritized for sprints.

84. JSON**Overview:**

JSON (JavaScript Object Notation) is a lightweight data-interchange format used for transmitting data between web servers and clients in a human-readable format.

Key Features:

- Lightweight and easy to parse
- Widely supported across programming languages
- Ideal for transmitting structured data

Example Use Case:

An API-driven application uses JSON to send data between the frontend and backend, exchanging information like user profiles and transaction history.

85. JSON API**Overview:**

The JSON API is a standard for building APIs in a consistent way, ensuring that requests and responses between clients and servers are structured in a predictable format.

Key Features:

- Consistent format for API requests and responses
- Supports pagination, filtering, and sorting
- Widely adopted for building RESTful APIs

Example Use Case:

A SaaS platform uses JSON API to build its backend, ensuring that API requests for fetching user data are structured in a standardized format for easy integration with third-party services.

86. JSON Fetcher**Overview:**

JSON Fetcher is a tool for fetching and processing JSON data from external APIs, simplifying the process of making HTTP requests and handling the responses.

Key Features:

- Simplifies HTTP requests to fetch JSON data
- Handles asynchronous data fetching
- Supports RESTful APIs and third-party integrations

Example Use Case:

A frontend application uses JSON Fetcher to make API calls to a weather service, fetching real-time weather data and displaying it on the UI.

87. Kafka

Overview:

Apache Kafka is a distributed streaming platform that allows for building real-time data pipelines and applications. It is designed to handle large volumes of real-time event data.

Key Features:

- Distributed and fault-tolerant architecture
- Real-time data streaming and processing
- High throughput for handling large event streams

Example Use Case:

A fintech company uses Kafka to stream real-time transaction data from their payment processing system to an analytics engine, providing instant insights into customer behavior.

88. kdb+**Overview:**

kdb+ is a high-performance time-series database optimized for financial data and used for real-time analytics, often in trading platforms and risk management systems.

Key Features:

- Optimized for time-series data
- High-performance data ingestion and querying
- Commonly used in financial services for trading and risk analysis

Example Use Case:

A trading firm uses kdb+ to store and analyze stock market data in real-time, helping them make split-second trading decisions based on historical trends and real-time signals.

89. Loki**Overview:**

Loki is a log aggregation system designed to work with Grafana, providing a cost-effective and scalable way to store, query, and visualize logs without the need for full-text indexing.

Key Features:

- Efficient log aggregation without full-text indexing
- Integration with Grafana for visualization
- Scalable architecture for large log volumes

Example Use Case:

A DevOps team uses Loki to store and visualize logs from their containerized applications, enabling them to monitor and debug issues in real-time.

90. Looker

Overview:

Looker is a business intelligence platform that allows organizations to explore, analyze, and visualize data in real-time, providing insights to inform decision-making.

Key Features:

- Real-time data exploration and analysis
- Customizable dashboards and visualizations
- Integration with databases and data lakes

Example Use Case:

A marketing team uses Looker to analyze customer data and sales performance, building custom dashboards to track key metrics and optimize marketing campaigns.

91. Materialize

Overview:

Materialize is a streaming SQL database that allows for real-time, incremental data updates and materialized views, enabling fast query performance on live data streams.

Key Features:

- Real-time data streaming with SQL queries
- Incremental updates to materialized views
- Fast query performance on live data

Example Use Case:

A financial analytics company uses Materialize to process real-time stock market data and provide instant insights to traders, ensuring that all queries reflect the latest market conditions.

92. MaxCompute

Overview:

MaxCompute is a fully managed big data processing platform provided by Alibaba Cloud, designed for large-scale data warehousing and batch processing.

Key Features:

- Scalable data warehousing for big data
- Supports SQL-like queries for batch processing
- Fully managed with security and compliance features

Example Use Case:

An eCommerce platform uses MaxCompute to process large volumes of sales data, generating reports and insights for product recommendations and inventory management.

93. Microsoft SQL Server

Overview:

Microsoft SQL Server is a relational database management system that provides enterprise-grade performance, security, and analytics for mission-critical applications.

Key Features:

- Relational database with SQL support
- High availability and disaster recovery
- Advanced analytics and machine learning integration

Example Use Case:

A financial institution uses SQL Server to manage their customer database, ensuring high availability and security for sensitive financial data.

94. Mock

Overview:

Mock is a tool used for creating mock data and APIs, allowing developers to simulate API responses and test their applications without needing a live backend.

Key Features:

- Create mock APIs and data
- Simulate various HTTP responses
- Useful for testing frontend applications

Example Use Case:

A frontend developer uses Mock to simulate API responses while building a new feature, allowing them to test the UI before the backend is fully implemented.

95. MonetDB

Overview:

MonetDB is an open-source columnar database designed for high-performance analytics and business intelligence workloads, providing fast query execution on large datasets.

Key Features:

- Columnar storage for fast analytics
- Optimized for business intelligence workloads
- Supports SQL and advanced analytics

Example Use Case:

A healthcare company uses MonetDB to analyze patient data from electronic health records, generating reports and insights to improve patient care.

96. MongoDB**Overview:**

MongoDB is a NoSQL database that stores data in flexible, JSON-like documents, making it ideal for applications with changing data schemas.

Key Features:

- Document-oriented NoSQL database
- Flexible schema design for dynamic applications
- Horizontal scalability for large datasets

Example Use Case:

A content management system (CMS) uses MongoDB to store user-generated content in a flexible format, allowing the system to adapt to different types of data without needing schema changes.

97. MQTT**Overview:**

MQTT (Message Queuing Telemetry Transport) is a lightweight messaging protocol designed for IoT devices, providing reliable data transmission over unreliable networks.

Key Features:

- Lightweight and efficient protocol
- Designed for IoT and M2M (Machine-to-Machine) communication
- Supports low-bandwidth, high-latency networks

Example Use Case:

A smart home system uses MQTT to transmit sensor data from devices like thermostats and motion detectors, enabling real-time control and monitoring of the home environment.

98. MySQL**Overview:**

MySQL is an open-source relational database management system that is widely used for web applications, offering fast performance and a large ecosystem of tools and support.

Key Features:

- Relational database with SQL support
- Fast and reliable performance

- Large ecosystem and community support

Example Use Case:

A web application uses MySQL to manage user accounts and transactional data, providing a stable and scalable solution for a growing user base.

99. Neo4j Datasource**Overview:**

Neo4j is a graph database that stores data in nodes and relationships, making it ideal for applications that require deep, complex queries on connected data, such as social networks.

Key Features:

- Graph-based data storage and querying
- Optimized for deep, complex relationships
- Cypher query language for graph data

Example Use Case:

A social media platform uses Neo4j to model and analyze user connections, enabling features like friend recommendations and community detection.

100. netdata**Overview:**

Netdata is a distributed, real-time monitoring and troubleshooting platform that collects thousands of metrics from servers, containers, and applications.

Key Features:

- Real-time performance monitoring
- Low-resource usage and lightweight agent
- Interactive visualizations and dashboards

Example Use Case:

A DevOps team uses netdata to monitor the performance of their production servers, receiving real-time alerts for high CPU usage or memory leaks.

101. Netlify**Overview:**

Netlify is a web development platform that automates builds, deployments, and scaling for modern web projects, particularly static websites and JAMstack applications.

Key Features:

- Automated CI/CD pipelines for web development

- Integrated CDN for fast, global content delivery
- Serverless functions for dynamic web applications

Example Use Case:

A web development team uses Netlify to deploy their static site automatically whenever changes are pushed to their Git repository, taking advantage of the integrated CDN for fast loading times.

102. New Relic**Overview:**

New Relic is a software analytics and performance monitoring platform that helps developers and DevOps teams track the health and performance of their applications and infrastructure.

Key Features:

- Real-time application performance monitoring (APM)
- Full-stack observability across applications, infrastructure, and logs
- Customizable dashboards and alerts

Example Use Case:

A DevOps team uses New Relic to monitor the performance of their microservices, identifying performance bottlenecks and improving the user experience of their application.

103. Node Graph API**Overview:**

Node Graph API is an API designed to allow developers to create and manipulate graph structures programmatically, providing tools to work with graph-based data structures.

Key Features:

- API for graph creation and manipulation
- Optimized for graph-based data
- Supports deep, complex queries on graph structures

Example Use Case:

A recommendation engine uses Node Graph API to build and query a graph of user interactions, providing personalized product recommendations based on relationships between users and items.

104. OData**Overview:**

OData (Open Data Protocol) is a standardized protocol for building and consuming RESTful APIs, allowing clients to query and update data in a consistent, uniform way.

Key Features:

- Standardized REST API protocol
- Supports querying, filtering, and pagination
- Broad adoption across Microsoft and SAP platforms

Example Use Case:

A business intelligence tool uses OData to query and update data in a CRM system, providing real-time insights into sales performance.

105. openGemini**Overview:**

OpenGemini is an open-source time-series database designed for high-speed data ingestion and low-latency queries, making it ideal for monitoring and IoT applications.

Key Features:

- High-speed data ingestion for time-series data
- Low-latency queries for real-time analysis
- Optimized for monitoring and IoT use cases

Example Use Case:

A smart manufacturing system uses openGemini to monitor production line sensors in real-time, allowing operators to quickly identify and address equipment malfunctions.

106. openHistorian**Overview:**

openHistorian is an open-source data historian designed for storing and retrieving time-series data, often used in industrial applications to track sensor data from machinery.

Key Features:

- Time-series data storage and retrieval
- Designed for high-frequency data collection
- Integration with SCADA and industrial systems

Example Use Case:

A factory uses openHistorian to collect and store time-series data from their machines, allowing them to analyze equipment performance and schedule preventive maintenance.

107. OpenSearch**Overview:**

OpenSearch is an open-source search and analytics engine derived from Elasticsearch, used for full-text search, log analytics, and real-time data visualization.

Key Features:

- Full-text search and analytics engine
- Real-time data visualization and log analytics
- Integration with Kibana for dashboards

Example Use Case:

An eCommerce platform uses OpenSearch to power its search functionality, providing users with fast, accurate product search results.

108. OpenTSDB**Overview:**

OpenTSDB is a scalable, distributed time-series database designed for storing and analyzing large amounts of time-series data, often used in monitoring and IoT applications.

Key Features:

- Distributed time-series database
- Optimized for high-frequency data collection
- Supports long-term data storage

Example Use Case:

A telecommunications company uses OpenTSDB to store and analyze network performance metrics, helping them identify and resolve potential issues before they impact customers.

109. Optimiz-SeqOne Plugin**Overview:**

Optimiz-SeqOne is a network performance monitoring plugin that allows for the collection and analysis of network traffic data, providing real-time insights into network health.

Key Features:

- Real-time network performance monitoring
- Traffic data collection and analysis
- Integration with other network management tools

Example Use Case:

A telecommunications company uses Optimiz-SeqOne to monitor network traffic across their infrastructure, identifying potential bottlenecks and ensuring optimal performance for their customers.

110. Oracle Cloud Infrastructure Logs

Overview:

Oracle Cloud Infrastructure Logs provides a fully managed log management service, allowing organizations to collect, store, and analyze log data from Oracle Cloud resources.

Key Features:

- Fully managed log management service
- Real-time log collection and analysis
- Integration with Oracle Cloud services

Example Use Case:

A financial services company uses Oracle Cloud Infrastructure Logs to monitor and analyze log data from their cloud infrastructure, ensuring compliance with regulatory requirements and identifying potential security threats.

111. Oracle Cloud Infrastructure Metrics

Overview:

Oracle Cloud Infrastructure Metrics provides real-time metrics monitoring for Oracle Cloud resources, allowing users to track performance, set alerts, and create custom dashboards.

Key Features:

- Real-time performance monitoring for Oracle Cloud
- Customizable dashboards and alerts
- Integration with Oracle Cloud services

Example Use Case:

An eCommerce platform uses Oracle Cloud Infrastructure Metrics to monitor the performance of their cloud-based infrastructure, ensuring optimal uptime and responsiveness for customers.

112. Oracle Database

Overview:

Oracle Database is a multi-model database management system designed for mission-critical applications, providing high availability, security, and scalability for enterprise workloads.

Key Features:

- Relational database with SQL support
- Advanced analytics and machine learning integration
- High availability and disaster recovery

Example Use Case:

A financial institution uses Oracle Database to manage their transactional data, ensuring high availability and security for critical business operations.

113. Orbit

Overview:

Orbit is a platform for building and managing online communities, helping organizations track and analyze user engagement, interactions, and growth metrics.

Key Features:

- Community management and growth tracking
- User engagement analytics
- Integration with communication tools like Slack and Discord

Example Use Case:

A SaaS company uses Orbit to track engagement in their online community, analyzing user interactions and identifying power users who contribute to product discussions.

114. OSIsoft-PI

Overview:

OSIsoft PI is a data historian platform that collects, stores, and analyzes real-time operational data from industrial processes, providing insights for optimization and predictive maintenance.

Key Features:

- Real-time data collection from industrial processes
- Time-series data storage and analytics
- Integration with SCADA, DCS, and other industrial systems

Example Use Case:

An oil and gas company uses OSIsoft PI to monitor equipment performance at their refineries, helping them predict and prevent equipment failures and optimize production.

115. PagerDuty

Overview:

PagerDuty is an incident management platform that helps DevOps and IT teams detect, escalate, and resolve incidents quickly, minimizing downtime and improving system reliability.

Key Features:

- Automated incident detection and escalation
- On-call scheduling and notification management
- Real-time incident tracking and resolution

Example Use Case:

A DevOps team uses PagerDuty to manage on-call schedules and automatically escalate incidents when a critical system goes down, ensuring fast response times and minimal downtime.

116. Parca

Overview:

Parca is an open-source continuous profiling platform for analyzing CPU and memory usage in applications over time, helping developers optimize performance and resource consumption.

Key Features:

- Continuous profiling of CPU and memory usage
- Real-time performance analysis
- Integration with Kubernetes and other orchestration tools

Example Use Case:

A cloud-native company uses Parca to monitor and optimize the performance of their Kubernetes workloads, identifying inefficient code and reducing resource consumption.

117. Parseable

Overview:

Parseable is a log management and analytics platform that enables organizations to collect, store, and search log data from multiple sources, providing insights for troubleshooting and performance optimization.

Key Features:

- Log data collection and search
- Real-time analytics and visualization
- Integration with monitoring and observability tools

Example Use Case:

An IT operations team uses Parseable to collect and analyze log data from their servers, helping them troubleshoot issues and optimize system performance.

118. Pixie Grafana Datasource Plugin

Overview:

The Pixie Grafana Datasource Plugin allows users to visualize Pixie-generated metrics and traces within Grafana, providing insights into application performance and behavior in Kubernetes environments.

Key Features:

- Visualize Pixie metrics and traces in Grafana
- Real-time monitoring and observability for Kubernetes
- Customizable dashboards and alerts

Example Use Case:

A DevOps team uses the Pixie Grafana Datasource Plugin to monitor the performance of their Kubernetes-based microservices, identifying performance bottlenecks and optimizing resource usage.

119. PNP**Overview:**

PNP is a graphing and monitoring tool designed for integrating performance data from Nagios, allowing users to visualize and analyze performance metrics over time.

Key Features:

- Graphing and monitoring for Nagios performance data
- Customizable dashboards for visualizing metrics
- Integration with Nagios and other monitoring tools

Example Use Case:

A systems administrator uses PNP to visualize server performance metrics collected by Nagios, helping them track CPU usage, disk I/O, and network traffic over time.

120. PostgreSQL**Overview:**

PostgreSQL is an open-source relational database management system known for its advanced features, extensibility, and SQL compliance. It is widely used for applications requiring complex queries and data integrity.

Key Features:

- Advanced SQL compliance and extensibility
- Support for ACID transactions
- JSON and JSONB support for NoSQL-like features

Example Use Case:

A web application uses PostgreSQL to manage user accounts and transactional data, benefiting from its SQL compliance and support for complex queries.

121. Prometheus**Overview:**

Prometheus is an open-source monitoring and alerting toolkit designed for reliability and scalability. It collects metrics from systems and applications, storing them as time-series data.

Key Features:

- Time-series data collection and storage

- Powerful query language (PromQL) for metrics analysis
- Integration with Grafana for visualization

Example Use Case:

A cloud-based application uses Prometheus to monitor system performance and set up alerts for high CPU usage, ensuring the application remains responsive under load.

122. Prometheus AlertManager Datasource

Overview:

The Prometheus AlertManager Datasource enables Grafana to visualize and manage alerts generated by Prometheus AlertManager, helping DevOps teams track and respond to incidents.

Key Features:

- Visualization of Prometheus alerts in Grafana
- Customizable alert rules and thresholds
- Real-time incident tracking and management

Example Use Case:

A DevOps team uses the Prometheus AlertManager Datasource in Grafana to track alerts for their containerized applications, ensuring critical incidents are addressed promptly.

123. Propel

Overview:

Propel is a continuous deployment platform designed for modern applications, providing automated testing, deployment, and rollback functionality for faster and more reliable releases.

Key Features:

- Automated testing and deployment pipelines
- Continuous delivery and rollback support
- Integration with CI/CD tools

Example Use Case:

A software development team uses Propel to automate the deployment of their microservices, reducing the time and effort required to release new features and bug fixes.

124. Proton

Overview:

Proton is a cloud infrastructure automation tool that simplifies the deployment of containerized applications and microservices, providing templates for infrastructure and application management.

Key Features:

- Infrastructure automation for containerized applications
- Predefined templates for application management
- Integration with CI/CD tools for automated deployment

Example Use Case:

A DevOps team uses Proton to manage the deployment of their containerized microservices, automating the provisioning of infrastructure and ensuring consistent deployments across environments.

125. QuestDB**Overview:**

QuestDB is an open-source time-series database optimized for high-performance data ingestion and querying, making it ideal for financial and IoT applications that require real-time analysis.

Key Features:

- High-performance time-series database
- Real-time data ingestion and querying
- SQL support for complex queries

Example Use Case:

A trading platform uses QuestDB to store and analyze real-time stock market data, providing traders with up-to-the-second insights and alerts.

126. Quickwit**Overview:**

Quickwit is a distributed search engine optimized for fast, scalable log and metric analysis, providing real-time insights into system and application performance.

Key Features:

- Distributed architecture for scalable log analysis
- Fast search and retrieval of log and metric data
- Integration with monitoring and observability tools

Example Use Case:

A cloud-native company uses Quickwit to search and analyze logs from their distributed applications, enabling them to quickly identify and resolve performance issues.

127. Rabbitmq**Overview:**

RabbitMQ is an open-source message broker that allows applications to communicate

asynchronously by sending and receiving messages through queues, enabling distributed systems to scale effectively.

Key Features:

- Asynchronous messaging with queues
- Supports multiple messaging protocols (AMQP, MQTT)
- Highly scalable for distributed applications

Example Use Case:

A microservices-based application uses RabbitMQ to communicate between services, ensuring messages are delivered reliably even when some services are temporarily unavailable.

128. Redis

Overview:

Redis is an open-source, in-memory data store used as a database, cache, and message broker. It supports key-value storage and offers fast performance for real-time applications.

Key Features:

- In-memory key-value store with fast access
- Support for data structures like lists, sets, and hashes
- High availability with replication and clustering

Example Use Case:

A real-time analytics platform uses Redis to store session data and user activity, providing low-latency access for analytics and reporting.

129. Riemann streams

Overview:

Riemann is a real-time event stream processing system designed for monitoring distributed systems, aggregating events, and triggering alerts based on custom rules.

Key Features:

- Real-time event stream processing
- Customizable event aggregation and alerting
- Integration with monitoring and alerting tools

Example Use Case:

A cloud-native application uses Riemann to process and aggregate metrics from their infrastructure in real-time, triggering alerts for high CPU usage and service failures.

130. RunReveal

Overview:

RunReveal is a real-time monitoring and alerting platform for cloud-native applications, providing insights into performance, resource usage, and system health.

Key Features:

- Real-time monitoring and alerting for cloud applications
- Insights into system performance and resource usage
- Customizable dashboards and alerts

Example Use Case:

A DevOps team uses RunReveal to monitor the performance of their Kubernetes clusters, ensuring that resource usage is optimized and any performance issues are addressed promptly.

131. Salesforce**Overview:**

Salesforce is a cloud-based customer relationship management (CRM) platform that helps businesses manage customer interactions, sales, and support services.

Key Features:

- CRM platform for managing customer data and interactions
- Customizable workflows for sales and support processes
- Integration with third-party tools and services

Example Use Case:

A sales team uses Salesforce to track customer leads and sales opportunities, automating follow-ups and ensuring that no opportunities are missed.

132. SAP HANA®**Overview:**

SAP HANA is an in-memory, column-oriented database management system designed for high-performance transaction processing and real-time analytics in enterprise environments.

Key Features:

- In-memory processing for real-time analytics
- Column-oriented storage for optimized query performance
- Integrated with SAP applications for enterprise use cases

Example Use Case:

A global enterprise uses SAP HANA to manage their financial data, enabling real-time reporting and analytics across their global operations.

133. Scanner

Overview:

Scanner is a tool for automated vulnerability scanning in applications, systems, and networks, helping organizations identify and fix security issues before they can be exploited.

Key Features:

- Automated vulnerability scanning
- Detailed reporting on security risks
- Integration with CI/CD pipelines for continuous security testing

Example Use Case:

A security-conscious development team uses Scanner to automatically scan their applications for vulnerabilities during the build process, ensuring that any security issues are identified and fixed before deployment.

134. Sentry

Overview:

Sentry is an error tracking and performance monitoring tool that helps developers monitor and fix crashes and bugs in real-time, improving application stability and user experience.

Key Features:

- Real-time error tracking and reporting
- Performance monitoring for applications
- Integration with popular programming languages and frameworks

Example Use Case:

A mobile app development team uses Sentry to track crashes and performance issues in their app, allowing them to fix bugs quickly and improve the user experience.

135. ServiceNow

Overview:

ServiceNow is a cloud-based platform for IT service management (ITSM), providing tools for managing incidents, requests, and changes within IT departments.

Key Features:

- ITSM platform for managing incidents, requests, and changes
- Automation of workflows for IT processes
- Customizable dashboards and reporting

Example Use Case:

An IT department uses ServiceNow to track and manage incidents, ensuring that any service disruptions are resolved quickly and efficiently.

136. ServiceNow Cloud Observability

Overview:

ServiceNow Cloud Observability is a monitoring and observability solution that helps organizations track the performance and health of their cloud infrastructure and applications.

Key Features:

- Real-time monitoring of cloud infrastructure
- Integration with other ServiceNow ITSM tools
- Customizable dashboards and alerts

Example Use Case:

A cloud-native company uses ServiceNow Cloud Observability to monitor their AWS infrastructure, ensuring that any performance issues are detected and resolved before they impact users.

137. Shoreline

Overview:

Shoreline is an incident automation platform designed for DevOps and SRE teams, allowing them to automate common operational tasks and respond to incidents more quickly.

Key Features:

- Automation of common operational tasks
- Real-time incident response and remediation
- Integration with monitoring and alerting tools

Example Use Case:

A DevOps team uses Shoreline to automate the process of restarting services when they become unresponsive, reducing downtime and improving system reliability.

138. simple grpc datasource

Overview:

The simple gRPC datasource plugin for Grafana enables users to collect and visualize data from gRPC-based services, providing real-time insights into application performance.

Key Features:

- Collect and visualize data from gRPC services
- Real-time monitoring and alerting
- Integration with Grafana dashboards

Example Use Case:

A microservices-based application uses the simple gRPC datasource plugin to monitor the performance of its gRPC services, ensuring that latency and errors are kept to a minimum.

139. SkyWalking**Overview:**

SkyWalking is an open-source observability platform designed for monitoring and analyzing distributed systems, providing tracing, metrics, and logging for microservices.

Key Features:

- Distributed tracing for microservices
- Metrics and logging for performance analysis
- Real-time monitoring and alerting

Example Use Case:

A cloud-native application uses SkyWalking to trace and monitor requests across its microservices architecture, identifying bottlenecks and improving performance.

140. Snowflake**Overview:**

Snowflake is a fully managed cloud data warehouse that allows organizations to store, analyze, and share large volumes of structured and semi-structured data in real-time.

Key Features:

- Cloud-based data warehousing and analytics
- Support for structured and semi-structured data
- Scalable architecture for large datasets

Example Use Case:

A retail company uses Snowflake to analyze sales data from multiple regions, generating real-time reports and insights to optimize inventory and marketing strategies.

141. SPARQL**Overview:**

SPARQL (SPARQL Protocol and RDF Query Language) is a query language used to retrieve and manipulate data stored in Resource Description Framework (RDF) format, commonly used in semantic web applications.

Key Features:

- Query and manipulate RDF data
- Designed for semantic web and linked data applications

- Integration with ontologies and knowledge graphs

Example Use Case:

A knowledge management platform uses SPARQL to query a large knowledge graph, retrieving relevant information based on complex relationships between entities.

142. Splunk**Overview:**

Splunk is a data platform designed for searching, monitoring, and analyzing machine-generated data, providing real-time insights into infrastructure, applications, and security.

Key Features:

- Search and analyze machine-generated data in real-time
- Dashboards and visualizations for log data
- Integration with security, observability, and IT operations tools

Example Use Case:

A cybersecurity team uses Splunk to monitor log data from firewalls, servers, and applications, detecting potential security threats and responding to incidents in real-time.

143. Splunk Infrastructure Monitoring**Overview:**

Splunk Infrastructure Monitoring provides real-time visibility into the performance and health of infrastructure and applications, helping organizations monitor and optimize resource usage.

Key Features:

- Real-time infrastructure and application monitoring
- Alerts and dashboards for performance analysis
- Integration with Splunk's observability platform

Example Use Case:

A cloud-native company uses Splunk Infrastructure Monitoring to track the performance of their AWS infrastructure, ensuring that resource usage is optimized and any performance issues are resolved quickly.

144. SQLite**Overview:**

SQLite is a lightweight, serverless, self-contained relational database engine commonly used in mobile and embedded systems, providing fast, simple database management for applications.

Key Features:

- Lightweight, serverless database engine

- Full SQL support for querying and transactions
- Ideal for mobile and embedded applications

Example Use Case:

A mobile app uses SQLite to store user data and preferences locally, enabling offline access to the app's features and seamless data synchronization when the user reconnects to the internet.

145. Sqlyze Datasource**Overview:**

Sqlyze is a data visualization and analysis platform designed for SQL databases, providing real-time query execution and insights into relational data.

Key Features:

- SQL query execution and visualization
- Real-time insights into relational databases
- Integration with popular SQL databases

Example Use Case:

A data analyst uses Sqlyze to query and visualize customer data stored in a MySQL database, generating reports on customer behavior and sales trends.

146. Strava**Overview:**

Strava is a social fitness network that allows users to track and analyze their physical activities, such as running and cycling, by recording GPS data and providing insights into performance.

Key Features:

- GPS tracking of physical activities
- Performance analysis and insights
- Integration with fitness devices and social features

Example Use Case:

An athlete uses Strava to track their daily runs, analyzing their pace, distance, and heart rate to improve performance over time.

147. Streamr**Overview:**

Streamr is a decentralized platform for real-time data streaming, enabling users to publish, subscribe to, and monetize data streams without relying on centralized platforms.

Key Features:

- Decentralized real-time data streaming
- Publish and subscribe to data streams
- Monetize data through the Streamr network

Example Use Case:

A smart city project uses Streamr to publish real-time data from traffic sensors, enabling third-party developers to subscribe to the data and create traffic management applications.

148. Sumo Logic**Overview:**

Sumo Logic is a cloud-native log management and analytics platform designed for DevOps, IT, and security teams, providing real-time insights into infrastructure, applications, and security.

Key Features:

- Real-time log management and analytics
- Integration with DevOps, IT, and security tools
- Customizable dashboards and alerts

Example Use Case:

A DevOps team uses Sumo Logic to monitor log data from their containerized applications, enabling them to detect performance issues and optimize resource usage in real-time.

149. Sun and Moon**Overview:**

Sun and Moon is a data source that provides real-time information about solar and lunar events, such as sunrise, sunset, moon phases, and lunar eclipses.

Key Features:

- Real-time solar and lunar event tracking
- Data on sunrise, sunset, moon phases, and eclipses
- Integration with weather and astronomical tools

Example Use Case:

An outdoor event planning company uses Sun and Moon data to schedule events based on sunset times and moon phases, ensuring that events are held at the most visually appealing times.

150. SurrealDB**Overview:**

SurrealDB is a multi-model, distributed database that supports both SQL and NoSQL querying, enabling developers to work with structured and semi-structured data in a single platform.

Key Features:

- Multi-model database with SQL and NoSQL support
- Distributed architecture for scalability
- Flexible querying for structured and semi-structured data

Example Use Case:

A social media platform uses SurrealDB to store user-generated content and relationships between users, taking advantage of its multi-model support for both structured and unstructured data.

151. TDengine Datasource**Overview:**

TDengine is an open-source time-series database optimized for IoT, big data, and industrial applications, providing fast ingestion and querying of high-volume time-series data.

Key Features:

- High-performance time-series data storage
- Optimized for IoT and big data applications
- Fast data ingestion and querying

Example Use Case:

An IoT company uses TDengine to collect and store sensor data from connected devices, providing real-time insights into device performance and environmental conditions.

152. Tempo**Overview:**

Tempo is a distributed tracing system designed for monitoring and troubleshooting microservices-based applications, providing real-time insights into request flows and performance bottlenecks.

Key Features:

- Distributed tracing for microservices
- Real-time monitoring and performance analysis
- Integration with Grafana for visualization

Example Use Case:

A cloud-native application uses Tempo to trace requests across their microservices architecture, identifying performance bottlenecks and optimizing the user experience.

153. TestData

Overview:

TestData is a tool for generating mock test data for applications, helping developers test their systems with realistic data without relying on production data.

Key Features:

- Generate realistic mock test data
- Supports multiple data types and formats
- Integration with testing frameworks

Example Use Case:

A QA team uses TestData to generate mock customer data for testing their eCommerce platform, ensuring that the application can handle various edge cases and data scenarios.

154. Thruk**Overview:**

Thruk is a monitoring web interface designed to integrate with Nagios, Icinga, and other monitoring tools, providing a centralized view of performance metrics and alerts.

Key Features:

- Centralized monitoring interface for Nagios and Icinga
- Real-time performance metrics and alerts
- Customizable dashboards and reporting

Example Use Case:

A systems administrator uses Thruk to monitor the performance of their server infrastructure, tracking CPU usage, memory, and network traffic from a single interface.

155. Tokio Console**Overview:**

Tokio Console is a runtime debugger for the Tokio asynchronous runtime in Rust, providing insights into the execution of asynchronous tasks and performance bottlenecks.

Key Features:

- Debugging and profiling for asynchronous Rust tasks
- Real-time insights into task execution and performance
- Integration with the Tokio runtime

Example Use Case:

A Rust developer uses Tokio Console to debug and optimize the performance of their asynchronous application, identifying bottlenecks and reducing latency.

156. Trino

Overview:

Trino (formerly PrestoSQL) is a distributed SQL query engine designed for running fast analytics queries on large datasets across multiple data sources, including Hadoop, AWS S3, and relational databases.

Key Features:

- Distributed SQL query engine for big data
- Query data from multiple sources (e.g., S3, Hadoop)
- High-performance, low-latency queries

Example Use Case:

A data analytics company uses Trino to query data stored in both AWS S3 and a MySQL database, providing fast, real-time insights across different data sources.

157. Vertica

Overview:

Vertica is a columnar database designed for high-performance analytics on large datasets, providing fast query execution and advanced analytics features for business intelligence and data science.

Key Features:

- Columnar storage for optimized analytics
- High-performance query execution
- Integration with data science and BI tools

Example Use Case:

A telecommunications company uses Vertica to analyze call data records and network performance metrics, generating reports and insights for optimizing their services.

158. Wavefront

Overview:

Wavefront is a real-time analytics and monitoring platform for cloud-native applications, providing insights into application performance, resource usage, and system health.

Key Features:

- Real-time monitoring and analytics
- Customizable dashboards and alerts
- Integration with cloud-native applications and infrastructure

Example Use Case:

A SaaS company uses Wavefront to monitor the performance of their AWS infrastructure, ensuring that resource usage is optimized and any performance issues are detected and resolved quickly.

159. WebSocket API

Overview:

WebSocket API is a protocol for full-duplex communication between clients and servers, enabling real-time data transmission for applications like chat systems and live updates.

Key Features:

- Full-duplex communication between clients and servers
- Real-time data transmission with low latency
- Widely used in chat, gaming, and live update applications

Example Use Case:

A real-time chat application uses WebSocket API to enable instant communication between users, ensuring that messages are transmitted in real-time with minimal latency.

160. Wild GraphQL Data Source

Overview:

Wild GraphQL Data Source is a plugin for Grafana that allows users to query and visualize data from GraphQL APIs, providing insights into application performance and behavior.

Key Features:

- Query and visualize data from GraphQL APIs
- Integration with Grafana for real-time monitoring
- Customizable dashboards and alerts

Example Use Case:

A SaaS company uses the Wild GraphQL Data Source to monitor data from their GraphQL API, visualizing performance metrics and detecting potential issues in real-time.

161. X-Ray

Overview:

AWS X-Ray is a distributed tracing system that helps developers analyze and debug microservices-based applications, providing end-to-end visibility into request flows and latency.

Key Features:

- Distributed tracing for microservices
- Real-time analysis of request flows and performance bottlenecks
- Integration with AWS services

Example Use Case:

A cloud-native application uses AWS X-Ray to trace requests across its microservices architecture, identifying performance bottlenecks and optimizing the user experience.

162. YDB**Overview:**

YDB is a distributed, horizontally scalable database service designed for high-performance applications that require low-latency access to large datasets, often used in financial services and IoT.

Key Features:

- Distributed, horizontally scalable architecture
- Low-latency access to large datasets
- Optimized for high-performance applications

Example Use Case:

A financial services company uses YDB to store and access real-time transactional data, ensuring that all queries are processed with low latency and high availability.

163. Yugabyte**Overview:**

Yugabyte is an open-source, distributed SQL database designed for high-performance, cloud-native applications, offering strong consistency and horizontal scalability.

Key Features:

- Distributed SQL database with strong consistency
- Horizontal scalability for cloud-native applications
- Support for both SQL and NoSQL workloads

Example Use Case:

A cloud-native eCommerce platform uses Yugabyte to manage its product catalog and customer data, ensuring fast, reliable access to data across multiple regions.

164. Zendesk**Overview:**

Zendesk is a customer service and engagement platform that helps businesses manage customer interactions across multiple channels, including email, chat, and social media.

Key Features:

- Customer service platform for managing interactions
- Multi-channel support for email, chat, social media, and more

- Integration with CRM and support tools

Example Use Case:

A support team uses Zendesk to manage customer inquiries and support requests, ensuring that all customer issues are tracked and resolved efficiently across multiple channels.

165. Zendesk Datasource Plugin

Overview:

The Zendesk Datasource Plugin allows users to query and visualize data from Zendesk within Grafana, providing insights into customer service performance and agent productivity.

Key Features:

- Query and visualize Zendesk data in Grafana
- Real-time insights into customer service metrics
- Customizable dashboards for agent performance and ticket resolution

Example Use Case:

A customer support manager uses the Zendesk Datasource Plugin to track ticket resolution times and agent productivity, identifying areas for improvement and optimizing support workflows.

166. Zipkin

Overview:

Zipkin is an open-source distributed tracing system that helps developers track and analyze the flow of requests across microservices, providing insights into performance bottlenecks and request latencies.

Key Features:

- Distributed tracing for microservices
- Real-time analysis of request flows and performance bottlenecks
- Integration with Kubernetes, Spring Cloud, and other platforms

Example Use Case:

A microservices-based application uses Zipkin to trace requests across its services, identifying bottlenecks and optimizing the overall performance of the system.

This concludes the detailed documentation for all 166 data sources. Each one is explained with an overview, key features, and a practical example use case to illustrate its real-world application. Let me know if you need further details on any specific data source!